

A Reevaluation of Why Crypto-detectors Fail: A Systematic Reevaluation of
Cryptographic Misuse Detection Techniques

Scott Marsden

Sherman Oaks, California, United States of America

Bachelor of Science, Elon University, 2020
Master of Science, College of William and Mary, 2023

A Dissertation presented to the Graduate Faculty
of the College of William & Mary in Candidacy for the Degree of
Master of Science

Department of Computer Science

College of William & Mary
April 2023

APPROVAL PAGE

This Dissertation is submitted in partial fulfillment of
the requirements for the degree of

Master's of Science in Computer Science

Scott Marsden

Approved by the Committee, April 2023

Committee Chair

Associate Professor Denys Poshyvanyk, Computer Science
College of William & Mary

Professor Dmitry Evt'yushkin, Computer Science
College of William & Mary

Professor Adwait Nadkarni, Computer Science
College of William & Mary

COMPLIANCE PAGE

Research approved by

Protection of Human Subjects Committee

Protocol number(s): PHSC-####-##-##-#####-protocol

PHSC-####-##-##-#####-protocol

Date(s) of approval: mm/dd/yyyy

mm/dd/yyyy

ABSTRACT

The correct use of cryptography is central to ensuring data security in modern software systems. Hence, several academic and commercial static analysis tools have been developed for detecting and mitigating crypto-API misuse. While developers are optimistically adopting these crypto-API misuse detectors (or crypto-detectors) in their software development cycles, this momentum must be accompanied by a rigorous understanding of their effectiveness at finding crypto-API misuse in practice. The original paper presents the MASC framework, which enables a systematic and data-driven evaluation of crypto-detectors using mutation testing. MASC was grounded in a comprehensive view of the problem space by developing a data-driven taxonomy of existing crypto-API misuse, containing 105 misuse cases organized among nine semantic clusters. 12 generalizable usage-based mutation operators were developed and three mutation scopes that can expressively instantiate thousands of compilable variants of the misuse cases for thoroughly evaluating crypto-detectors. Using MASC, nine major crypto-detectors were evaluated and 19 unique, undocumented flaws that severely impact the ability of crypto-detectors to discover misuses in practice were found.

For my thesis I built upon this previous research and greatly expanded the MASC framework. MASC was expanded in all areas by adding new functionality, new operators, new misuses, and expanding the taxonomy. In addition, I reevaluated the most up to date versions of the original 9 crypto-detectors and evaluated 5 additional crypto-detectors. On top of this I also doubled the amount of applications I used to evaluate the tools.

The new functionality that was added to MASC includes a tool for users to obtain mutants based on a type of static analysis sensitivities such as: path sensitivity, alias sensitivity, context sensitivity, object sensitivity, and flow sensitivity. I researched and defined all these sensitivities and categorized each of the existing and new operators to fit under these sensitivities. This was designed to help expand user usability of the tool. In addition I also added a tool to parse SARIF output (a type of output produced by crypto-detectors) and determine which mutants were caught. This tool was designed to help with analyzing crypto-detectors and speed up the analysis time. In the original paper all results were determined by hand.

I expanded functionality of existing operators so that they could be more flexible and have new features. I also added 7 brand new operators to the group of existing operators. These operators allow for MASC to handle a variety of new misuse cases and provide many more options for mutations. This allows for more options to evaluate crypto-detectors and determine identify their flaws.

Finally MASC built a taxonomy of known crypto api misuses up till 2019. This taxonomy was updated to include any misuses discovered between 2019 and 2022. This resulted in 19 new papers being added to the taxonomy with a new total of 55 papers. This expansion both reinforces the misuses that were previously identified in the taxonomy and added several new misuses to the taxonomy as well.

Finally MASC built a taxonomy of known crypto api misuses up till 2019. This taxonomy was updated to include any misuses discovered between 2019 and 2022. This resulted in 19 new papers being added to the taxonomy with a new total of 55 papers. This expansion both reinforces the misuses that were previously identified in the taxonomy and added several new misuses to the taxonomy as well.

To analyze crypto-detectors I looked at both the 9 crypto-detectors evaluated in the original work and 5 new crypto-detectors. For the original crypto-detectors I evaluated them with the updated MASC against their most up to date versions to determine if old flaws that have previously been fixed have a tendency to reappear. Both old and new crypto-detectors were evaluated with mutated Android and Java applications that were used in the original MASC paper, 15 newly mutated Android and Java applications, and minimal examples of cryptographic misuse.

TABLE OF CONTENTS

Acknowledgments	iii
Dedication	iv
List of Tables	v
List of Figures	vi
1 Introduction	2
1.1 Overview	3
1.1.1 Motivation and Background	5
1.1.2 Threat Model	7
2 The MASC Framework	9
2.1 Overview	9
2.1.1 Extended Taxonomy	10
2.1.2 Extended Operators	14
2.1.3 Threat Based Mutation Scopes	17
3 Implementation	19
3.0.1 New Features	21
3.0.2 Sensitivity Evaluator	21
3.0.3 Automated Evaluation	23
3.0.4 MASC Web	26

4	Evaluation and Methodology	27
5	Results and Findings	32
5.0.1	Analyzing Flaws	32
5.0.2	Exhaustive Results	34
6	Conclusion	35
6.0.1	Limitations	35
6.0.2	Discussion	37
6.0.3	Lessons Learned	39
6.0.4	Conclusion	39
A	Appendix A	40
A.1	Code Snippets	40

ACKNOWLEDGMENTS

I would like to first acknowledge Amit Ami for helping guide me through the thesis. We met nearly weekly to help ensure I was making progress and that it was consistent with the original work. Amit was a great help as both a mentor and someone who could double check my work.

I would also like to thank the developers of all the tools that were evaluated within the paper. Without these developers working to provide these tools I would not be able to find ways to improve them. I also appreciate that they were open to a discussion to find ways to improve the tools.

Finally, for early portions of the project I worked in collaboration with other students for class projects. I would like to thank them as well for contributing to the project.

Insert heartfelt dedication here...

LIST OF TABLES

LIST OF FIGURES

A Reevaluation of Why Crypto-detectors Fail: A Systematic
Reevaluation of Cryptographic Misuse Detection Techniques

Chapter 1

Introduction

Cryptography is essential in making sure data is confidential and secure in the modern world. Making sure cryptography primitives are used correctly, however, has always been a difficult problem to solve whether it be in critical systems like banking or the wide spread misuse of cryptographic API found in mobile and web apps. These misuses can lead to vulnerabilities that allow for data leaks and cause sensitive data to be leaked. To solve this problem researchers have designed tools that can be integrated into the development pipeline that catch these misuses and allow developers to catch the problems before they are released to the public. These tools are known as crypto-API misuse detectors also known as crypto-detectors. These crypto detectors are vital to helping keep development of applications secure. However, these crypto-detectors are not without their own faults. Crypto-detectors see widespread use across IDEs, corporate testing suites, by source control platforms (such as GitHub), and commercial use. Crypto-detectors are relied upon by many developers that deal with sensitive information, there is an expectation and an assumption that they are reliable. Developers utilize these crypto-detectors to help have peace of mind that their code is now more secure. However, a lot is still unknown about how effective they are at crypto-API misuse despite how eager developers are to adopt them. The MASC framework presented in the original paper tried to provide a solution to this problem by providing a framework to evaluate crypto-detectors beyond simple manually created benchmarks. As described in the

original paper [insert citation] MASC is "the first systematic, data-driven framework that leverage the well-founded approach of Mutation Analysis for evaluating Static Crypto-API misuse detectors." MASC is used by a user in a imilar manner to typical mutation analysis. MASC is capable of "mutating Android/Java apps by seeding them mutants" containing crypto-API misuse. The mutated apps can then be anaylzed by a crypto-dector and seeing which mutants are not located during the analysis reveals both design and implementation flaws in that crypto-dector. For my thesis we took the already existing MASC framework and extended it. This was done by adding a variety of new functionality and greatly expanding the functionality that already existed. We expanded MASC in both depth and width. The updated MASC framework was also used to reevaluate Crypto-Dectors that were evaluated by the original paper and evaluate 5 additional crypto-dectors. With the new additions to MASC and new evaluation we were able to determine: do misuses that were reported and fixed have a tendency to reappear in future builds, have crypto-dectors improved since the original paper, can MASC discover new flaws in crypto-dectors, what are the chracterisitcs of these flaws, and what is the impact of the flaws on the effectiveness of crytpo-dectors in practice?

1.1 Overview

The original MASC framework was built based on a Crypto-API Misuse Taxonomy. At the time is was the first comprehensive taxnonoy of crypto-API misuse cases containing 105 cases. These were found using a data driven process that systematially identified, studiad, and extacted misuse cases from both academic and industrial sources published from 1998-2018. This taxonomy provided the main building block for designing mutants that can emulate realistic misuses. For my thesis we expanded this taxonomy to cover papers from 2019-2022. The extended taxonomy includes ... more papers and ... more misuses. MASC also had several Crypto-Mutation Operators and Scopes. These different abstractions were designed to allow flexible in MASC to create a large variety of feasi-

ble mutants. The threat model MASC is based on consists of three types of conditions that crypto-dectors are likely to face. In addition, MASC also consisted of usage base mutation operators which were general operators that contain common characteristics that can leverage a diverse set of crypto APIs to create misuse cases that can be found within the taxonomy. These operators consist of two main categories: restrictive, where a developer can only instantiate certain objects by providing values from a predefined set such as `Cipher.getInstance(<parameter>)` and flexible which consists of operators with a large amount of extensibility such as developers can customize the hostname verification component of the SSL/TLS handshake by creating a class extending the `HostnameVerifier`, and overriding its `verify` method, with any content. The original set of operators consisted of 12 operators (6 restrictive and 6 flexible). For my thesis we focused on expanding the amount of restrictive operators since more are necessary to fully represent the taxonomy. I added ... number of restrictive operators to the total count based both on misuses that were not represented from the original taxonomy and to represent the extended taxonomy. I also added restrictive operators to further extend our T3 threat model of an evasive developer since the original work was intentionally conservative with what was considered evasive. After more research it was determined that new operators could be created to more evasive and still fair to crypto-dectors. MASC also includes 3 types of mutation scopes that allow for seeding of mutants variable fidelity to realistic API-use and threats. I also expanded MASC to contain several new features. These consist of adding an automated analysis tool and the sensitivity runner `<change name>`. The automated analysis provides MASC with a way to automatically evaluate a crypto-dectors output and determine which misuses were caught or not. It leverages SARIF files (an file type similar to JSON) to determine where misuses were. It is also designed to catch flaws with crypto-dectors based on output. It can test a variety of cases against a crypto-decor that range in complexity and will report if a simple case fails to be caught and stop running, since a more complex case cannot be feasibly caught if the simple version fails. Since all evaluation from the original paper was done by hand and was very time consuming to ensure accuracy. This tool was designed to

help researchers and users save time while evaluating crypto-dectors and ensure accuracy. The other main new feature introduced as a part of MASC is the sensitivity runner. This tool defined the different types of sensitivities common to crypto-dectors: flow, object, context, path, and alias sensitivity. Certain tools claim to be built with these sensitivities in mind. These were used to categorize all the operators contained under one or more of these categories. MASC can then be run and will generate mutants only for the specified sensitivity type. This allows for a lower barrier to entry for users and to allow users to test mutants specifically against the claims of a crypto-decor. For the evaluation of crypto-dectors I used the same methodology as the original paper but extended the number of operators and the number of applications. I evaluated 8 of the major crypto-dectors used in the original paper and also evaluated 5 additional crypto-dectors. The original paper also evaluated the crypto-dectors with a 15 different Android/Java applications and a group of minimal cases. All of these same applications and minimal cases were used again in evaluation with the addition of 16 new Android/Java applications that were mutated using the newly extended MASC. Additionally, our findings were disclosed to the designers/maintainers of the tools.

1.1.1 Motivation and Background

The motivation for extending MASC still comes from the same idea as the original motivation behind MASC, "Insecure use of cryptographic APIs is the second more common cause of software vulnerabilities after data leaks. Many developers are likely to use crypto-dectors in their development pipeline since they are not experts in the area but wish to catch vulnerabilities before releasing software. This means that the reliability of crypto-dectors and their ability to locate misuses directly impacts the security of the end user and their data. Evaluating crypto-dectors is an ongoing problem. New misuses are discovered frequently and it is important to have a tool that can keep up with new misuses. The idea behind expanding the MASC framework came from the need to keep up with new misuses and still be a reliable tool for evaluating crypto-dectors. The crypto-dectors evaluated in

the original paper needed to be reevaluated because bugs that get fixed have a tendency to reappear in future patches. This is a common occurrence in the field of software engineering.[citation] The newest versions of these tools needed to be reevaluated to see if they have made any improvements with catching misuses and ensure that flaws that were found have truly been fixed. The set of evaluated crypto-detectors was also expanded so that flaws can be reported to other widely used tools to help improve them as well. In addition since the original paper was published new tools have appeared and gained traction such as Amazon CodeGuru and SonarQube. I wanted to ensure MASC conducted an evaluation on as many crypto-detectors as possible. Since MASC's viability was proven in the original work it was important to extend and improve the work as well as extend its reach. The MASC framework, like many software engineering projects, is an ever evolving framework that is designed to be a reliable way to help evaluate the effectiveness of crypto-detectors. The motivating example for this new component of the research would be a scenario like the following. Consider Alice, a Java Developer who uses SonarQube as a part of her development pipeline, a known state of the art crypto-detector known for being able to detect security vulnerabilities found in software prior to release. Alice had reported that there was a flaw in the crypto-detector that allowed the following line of code to pass without detection: `Cipher cipher = Cipher . getInstance (" des ") ;` She was told that in the new version of the tool that this would be fixed. Her team upgrades to the latest version and sees that it now detects this misuse. Then some time goes by and another new version is released. Alice's team upgrades to this but one of her team members, Bob, puts the same line of code in the software unaware that it is a misuse. It is able to get through to production since the crypto-detector did not detect it even though Alice reported it previously and is under the belief that it is no longer present. In addition I continued using the motivation from the original paper of Alice being an unaware developer using: `Cipher cipher = Cipher . getInstance (" des ") ;` and it not being detected by the crypto-detector. DES is the common version of the misuse, however, Java supports both the uppercase and lowercase versions of this misuse. A crypto-detector would be expected to detect this common mistake.

1.1.2 Threat Model

The threat model remains the same between both this extension of work as well as the original. In the original paper they defend the scope of their evaluation by leveraging the documentation of popular crypto-detectors to understand how the position their tool and what they claim to be capable of. When evaluating the new tools we looked through their documentation as well and ensured that they made similar claims as the originally evaluated crypto-detectors to ensure the model was consistent across all evaluated crypto-detectors. As an example snyk's documentation claims "Empower developers to become quasi-security professionals with Snyk Code's comprehensive security tooling." Once again similarly Amazon Code Guru claims "Our security detectors use machine learning and automated reasoning to analyze data flow to perform whole-program inter-procedural analysis, across classes, methods, and files to detect hard-to-find security vulnerabilities." and "Java Crypto Library Best Practices help you check common Java cryptography libraries, such as Javax.Crypto.Cipher, to identify that they are initialized and called correctly" or SonarQube says "maximum protection with taint analysis." All 5 of the newly evaluated crypto-detectors make similar claims of security assurance. Just like the original 9 crypto-detectors that were evaluated and now reevaluated they need to be evaluated by a 3rd party to verify their claims and assure developers that they can truly claim what they say they can. For our threat model we continued off the same model used in the original paper. It assumed that there are some circumstances where they might be deployed in adversarial circumstances due to their claims of being useful for security audits and similar tasks. There are some circumstances where security audits are required and one of the parties involved is opposed to this such as verification for deploying an app in the Google Play store. To reiterate the threat model consists of 3 types of adversaries (T1-T3). This threat model is how components of MASC are designed and guide how evaluation was conducted: T1 Benign developer, accidental misuse - this model assumes the developer accidentally misuses crypto-API, but attempts to detect and address the vulnerabilities with the assistance

of crypto-detector. T2 Benign developer, harmful fix - This scenario shares some similarities to the first as it assumes the developer does not fully understand crypto-API but they are using a tool to help identify misuses. When a misuse gets flagged the developer attempts to fix it but introduces a new vulnerability in its place. T3 Evasive developer, harmful fix - This assumes the developer is trying to finish the task quickly or with low effort and is intentionally trying to evade the crypto-detector. When the alert from a detector is received the developer will do quick fix that results in potentially hiding the misuse rather than truly fixing it. Once again like the original design of MASC the conditions are meant to mimic how crypto-detectors have to operate in practice and the evaluation is designed based on what the crypto-detectors should be detecting. However, there is a gap between what should be and what is. This is why just like the design of the original MASC I kept all use cases in consideration. However, I did try to further capture the T3 developer compared to the original work. In the original work the T3 developer was initially designed to be somewhat conservative to ensure that cases were as close to reality as possible. The original developer were aware that they could go further but they wanted to truly ensure that all examples were accurate to cases that could be found. For this extension it was discovered that there are cases far more complex and malicious than was originally realized. So in the expansion of MASC some operators were added to include more complex T3 cases.

Chapter 2

The MASC Framework

2.1 Overview

The original work proposed a framework for Mutation-based Analysis of Static Cryptomise detection techniques (or MASC). [Add Fig 1] shows the design of the MASC framework as described in the original work. The framework remains the same and all the ideas from this original framework were intentionally consistent throughout my work as well. Cryptographic libraries contain many sets of API with a variety of potential misuse cases, this resulted in an incredibly large design space. Due to this the original work decided to start MASC by creating a "data-driven taxonomy crypto API misuse." This allows for the the misuse cases to be grounded based on what is seen in practice. This both allows for a more focused design space as well as helps justify the misuses MASC puts into practice. For this extension this taxonomy was expanded to include more cases than the prior work and updated with newer misuses and additional confirmation of previous misuses. Misuses had to be handled in a way that allowed for some flexibility since misuses can appear in a variety of ways. For example the original paper provides the example of DES can be provided as a variable in `Cipher.getInstance(<parameter>)` and can also be provided in lowercase. To be able to express both of these without hard coding them as examples MASC is designed by defining usage-characteristicsof cryptographic APIs. These

can then be leveraged to design "general, usage based mutation operators." These operators being designed in this way allows for a single operator to be able to handle a variety of misuses. This was made clearer when the taxonomy was expanded due to the fact that it was found that some of the operators previously designed for MASC were already capable of expressing these misuses with little to no effort. A substantial amount of work was put into extending operators to cover misuses that were not already represented in both the original and extended taxonomy. By instantiating mutants by applying the mutation operators to the cases found in the taxonomy, MASC injects, the mutants into Java/Android applications. How mutations are seeded is based on one of three scopes: the main scope, exhaustive scope, and similarity scope. These scopes are designed to emulate practical scenarios described in the threat model. Once mutated the apps can then be used to analyze a crypto-decor targeted for evaluation. Based on the results produced by the target it is possible to determine based on undetected mutants where design and implementation flaws exist. This is especially true in the case of reemerging cases.

2.1.1 Extended Taxonomy

In expanding MASC as a framework it was necessary to update the taxonomy to bring it up to date from when the original taxonomy was created in 2018. Since the taxonomy is serving as expansion rather than a replacement the same methodology was used from its initial creation to extend it. This was done to ensure the taxonomy as a whole was consistent and one whole work rather than a two different components. In order to locate academic and research papers related to Crypto-API misuse, it was first necessary to identify a set of search terms that would adequately narrow the search space. The authors of the original MASC paper had already identified a sufficiently narrowing set of search terms. Specifically, these terms are shown in Figure .

These search terms were used, as well as combinations of these search terms, as a foundation in conducting our searches for academic and industrial papers published between 2019 and 2022 identifying Crypto-API misuse cases. This was also done to make ensure

the methods used to collect papers were consistent with the methods used in the original paper.

With respect to academic papers, searches were conducted using the search terms above in Figure , as well as combinations of those terms on Google Scholar, IEEE Xplore & ACM Digital Library as the primary search engines (these were also the engines used when creating the taxonomy for the original paper). In addition, each citing reference in each of the papers concerning Crypto-API misuse found through searching these databases and search engines was also looked at. By doing this, it was possible to locate approximately 5-6 academic papers, concerning Crypto-API misuse, that were not readily visible through searches alone. This method helped expand the search domain for obtaining relevant academic papers.

For industrial sources, Google, Bing, and StackOverflow were relied upon for searching. While it was possible to locate some relevant documents, they were few and far between compared to the academic papers that we uncovered. In an effort to ensure completeness of the search space, when an industrial document relevant to Crypto-API misuse was located, the references it cited if any were also considered.

Additionally, a search was conducted through the Common Weakness Enumeration (CWE) from the MITRE Corporation . The CWE is a database of known hardware and software vulnerabilities, which are classified and categorized. Many crypto-dectors use CWE as a base for misuses that they detect. Notably, it includes known weaknesses that involve software using an API (or Crypto API) in a manner contrary to its intended use. We searched the CWE for new misuse cases not already identified in the original MASC Taxonomy or not already revealed in academic papers located from 2019 to 2022. Unfortunately, the CWE did not contain any Crypto-API misuses that were not already contained in the MASC Taxonomy or based on our searches of academic papers from 2019 to 2022.

To ensure that selected papers were only relevant to Crypto-API misuse in the MASC Taxonomy extension, the inclusion and exclusion criteria outlined in the original MASC

paper was followed. Specifically, to decide whether to consider a paper for further analysis, we used the inclusion criterion that the paper should discuss Crypto-API misuse or its detection. Exclusion criterion was used that the Crypto-API misuse described by the paper will be excluded if it does not relate to the Java programming environment or ecosystem, if the paper was published prior to 2019 or if the paper did not contain relevant information related to the subject matter of MASC.

In addition to these criteria, an additional exclusion criterion was included specifically with respect to the MASC Taxonomy extension. That is, if a paper discussed Crypto-API-related misuses, but did not identify specific misuse cases in its text, the paper was added to the general list of sources. However, these papers were not included in the final list of taxonomy sources from which misuse cases were extracted for the MASC Taxonomy extension.

After compiling the master list of misuse sources, a misuse case extraction was performed. We had 2 researchers independently identify and record misuse cases present within each of the sources. After this extraction process occurred, the two people met and had what was termed an “agreement/disagreement meeting”, as described in the original paper, in order to discuss their findings with respect to our extractions. This was done to ensure consistency across the different papers and ensure that each misuse case was found as well as confirmed. If the two members had any sort of disagreement over if something was a misuse case or not a third independent party was brought in to ensure correctness.

Misuses were extracted and organized with the same methodology as the original paper. The same clusters were used that were designed for the original paper to organize each of the misuses into categories. The clusters were created based on two differentiating criteria: "(1) the security goal/property represented by the misuse case (e.g., secrecy, integrity, non-repudiation) and (2) its level of abstraction within the communication/computing stack (e.g., confidentiality in general, or confidentiality with respect to SSL/TLS)"

In addition to this another misuse extraction was performed at a later date based on the misuses that were extracted. I went through each paper in the taxonomy and located each

misuse that was listed for each of these new papers. This was done to confirm the work that had already been done but ensure that it was thorough. This step was meant to be redundant and ensure the correctness of the taxonomy but resulted in some disagreements of misuses within 3 of the sources and a disagreement with one of the new sources in the taxonomy as a whole.

The approach for this additional step was for us to look at each individual paper. Based on the updated taxonomy that was produced we would go through the paper and look for each misuse that was specified to be contained within this paper. If the misuse was found it would be marked down and confirmed. If it was not found that misuse in the taxonomy would be flagged and would be brought up again with the original third party. If it was confirmed it would stay in the taxonomy if not it would be removed. Since the taxonomy is an extension I wanted to give extra cation to ensure not only correctness but consistency to ensure the taxonomy was one whole versus being two seperate parts.

By the end of this process 13 new misuse cases were identified and added to the MASC Taxonomy. In addition I was also able to further confirm many of the misuses that were already present in the taxonomy with the additional papers. The new misuse use cases include:

- Storing sensitive data in Java String (5 occurrences)
- Key reuse in stream ciphers (1 occurrence)
- Use of expired keys (1 occurrence)
- No clearPassword call after using PBEKeySpec (1 occurrence)
- Using AES-CTR (1 occurrence)
- Weak algorithm for password-based encryption - PBKDF1 (1 occurrence)
- HMAC for TLS with MD5 (1 occurrence)
- Using RC5* (1 occurrence)

- Using ARCFOUR (1 occurrence)
- PBEWithMD5AndDES (1 occurrence)
- Hashing Credentials - SHA-224 (1 occurrence)
- Reusing IVs & key pairs (1 occurrence)
- Manually changing hostname verifier (1 occurrence)

I believe that this is a solid basis for MASC as a whole and helped give ideas for expansion of operators to ensure that MASC is keeping up with the changing requirements. This taxonomy will likely have to be updated again in the future since security misuses come up frequently and the field is ever changing. I feel that this taxonomy is a good representation of cryptographic misuses through 2022.

2.1.2 Extended Operators

When designing mutation operators I continued building them based on the tradeoff of representing as many misuses cases as possible while also creating a number of operators that can still feasibly maintained. As the project grows this will continue to be a challenge. This is why it is important when operators were designed to make them as flexible to as many misuse cases as possible. This requires a lot of time and planning to ensure the design of the operators is maintainable. The other goal of operators as laid out in the original work is that they are not designed to exploit general soundness limitations such as dynamic and implicit calls. This is important to ensure that the results that are found are actual flaws versus something that a tool cannot be reasonably expected to compute. The operators were designed to be expressive of multiple misuse cases to allow for more coverage. Design of operators is also guided by the threat model described previously. When an operator is designed which threat model it covers is considered as well. I felt there was a lot of room to expand upon the evasive developer (T3) so most of the new operators focus on this threat model.

The MASC framework consists of two main types of operators: flexible and restrictive operators. These two types are based on the common characteristics that were found amongst misuses of crypto-API. The restrictive operators are operators where a developer can only instantiate certain objects by providing values from a predefined set such the method `Cipher.getInstance(<parameter>)` only accepts predefined configuration values for the parameter in String form. While other crypto-APIs allow significant amount of customizability and extendability resulting in the flexible type of operators. For this work I exclusively focused on extending the restrictive operators.

Due to the nature of restrictive operators having more limitations it leaves plenty of room for expansion. In addition, since the taxonomy was expanded it was possible to utilize some of the ideas that already existed in the taxonomy as well as newly discovered misuse case. Since these operators have limitations, there are many possibilities for expanding the restrictive operators. For the extension of MASC we have focused heavily on expanding the number of restrictive operators as well as the functionality of some of the restrictive operators that already existed.

While expanding the operators there was also more of a focus placed on the (??) cases. The initial paper was intentionally conservative with what was considered an evasive developer. This was due to the work being new and wanting to ensure that the operators were considered fair based on research conducted. After some time and further research I felt that it was possible to push further with this type of developer and after seeing further examples of evasive developers in real life cases. I wanted to ensure MASC is well rounded when simulating different types of developers and provide the most use cases possible for crypto detectors.

OP₁₃: *Iterative Method Chaining* – Similar to in MASC **OP₅** (Method Chains) this operator implements method chaining to hide the value. Where this operator differs from **OP₅** is that it can take a value specified by the user and create that many method calls. The method calls are then chained in succession. Every method call will have a safe value until the final call which transforms the value into an unsafe value. This was created

to test the limits of how far crypto detectors check method calls and can allow a user to determine where their fault point is. Just like **OP₅** this behavior would simulate a (??) developer.

OP₁₄: *Iterative Nested Conditionals* – This operator shares some similarities with the idea behind **OP₁₃: *b*** – ut instead of using method calls based on the iteration value it creates nested conditionals based on the value. Within the if statement there is an unsafe value and in the else a safe value. The nested conditionals will always pass the unsafe value but like **OP₁₃: *t*** – his operator tests how many nested conditionals a crypto detector can evaluate.

OP₁₅: *Method Builder* – This case takes method calls to build the String of a vulnerable call. The object class has methods that each contain a letter such as “D”, “E”, and “S”. Then it has a method that will add the letters together by calling each method and setting them equal to the variable in the class. This variable would now be "DES" and could be passed into an unsafe method. This operator also simulates (??) behavior and is not something a developer would accidentally do.

OP₁₆: *Object Sensitive* – The operator creates two versions of a base object that has method calls to make a String either secure or insecure. One object sets the variable to secure while the other sets itself to insecure. Then we set the object that is currently secure equal to the object that contains the insecure String. Then the originally secure object String is passed into the vulnerable API. This is done to test how well crypto detectors can handle object sensitivity. This was designed to give MASC more options for testing object sensitivity since this was an area that was originally lacking. This behaviour also fits under our T3 developer.

OP₁₇: *Build Variable* – For this operator I converted an insecure String into a char array. Then when the vulnerability is passed into the API the char array is then converted back to a String during the method call. This would fall under the (??) developer. This is because a developer could be doing some form of String manipulation that requires converting a String to a char array and then passing that into the method call.

OP₁₈: *Substring* – The idea behind this case is a user pulling the misuse out of a substring. In this case we took something like “HelloWorldDES” and passed the substring “DES” into the vulnerable api. Once again this is a (??) developer since they might pull something out of String and pass it in not knowing that it is vulnerable. Since it is not being passed in simply as a String if the crypto detector does not process the change this is easily a miuse that could slip by.

OP₁₉: *Static Keystore* – This operator is designed to handle static bytes being passed in insecure ways. For example, if a developer attempted to pass static bytes that are stored in a variable into Android Keystore since this is considered unsecure behavior. This would emulate a T1 developer since this is a very easy mistake to make if someone is unaware of the rules associated with crypto API and vulnerabilities.

2.1.3 Threat Based Mutation Scopes

The mutation scopes are designed to emulate the three types of developers described in the threat model. The scopes try to closely simulate placements of vulnerable code for the benign (T1, T2) and evasive developers: The main scope creates base case mutations and generates simple Java files to be tested. This is used mainly to determine if a crypto-dector is even capable of detecting misuses in the most plain cases. These mutants are seeded in simple Java or Android templates at the begining of the main method. This ensures that the mutants are found and analyzed. This is meant to simulate the behavior of a beniegn (T1 and T2) developer. The exhaustive scope looks at seeding mutations in every possible location such as class definitions, conditional segments, method bodies and anonymous inner class object declarations. Note that it is seeded in places that ensure the app is still compilable. Typically this is used with a single miuse that is known to be detected by the operator and is used to determine how thorough the crypto-dectors when analyzing files. This scope is meant to analyze a T3 developer they may hide miuses in places that one would not normally expect misuses to appear. The similarity scope seeds mutants in places where a similar security API is already in use. It emulates taking possibly secure

uses and making them insecure while not overwriting the previous code. This emulates where the benign developers (T1 and T2) would potentially place misuses and is meant to test how well the crypto-dectors analyze realistic areas misuses might appear.

Chapter 3

Implementation

The implementation of MASC remains consist with the original work. It involves the same three original components: "(1) selecting misuse cases from the taxonomy for mutation, (2) implementing mutation operators that instantiate the miusse cases, and (3) seeding/inserting the instantiated mutants in Java/Android source code at targeted locations." The extension of the framework was build upon the foundation that was previously laid. All additions remained consistent with the implementation of the original work.

1. Selecting misuse cases from the Taxonomy: In the original work 19 misues were chosen from the taxonomy for mutation utilizing the original 12 operators. Misuses were chosen based on two main factors to ensure that different categories of cryptographic misuse were represented as well as ensuring the more prevelant cases appeared as well. For the extened work only a few additional misuse cases were added such as a static key in android keystore. However, with the expansion of the operators it possible to create many more mutants with the misuses that were already present. For this extension the previous misuses were leveraged when creating mutations. However, they were designed with other possible misuses in mind similarly to the original 12 operators.
2. Implenting mutants: The mutation operators described in the previous chapter along with the original 12 operators were desgined to be applied to one or more crypto-API, for instatiating specific misuse cases. The goal of generating mutants in programs is to ensure that they are still compilable. To ensure that

this remains possible MASC considered the necessary syntactic requirements of each API that is being implemented (such as the requirement of surrounding try-catch block with appropriate exception handling) and the semantic requirement of the specific misuse that is being instantiated. MASC used Hava Reflection to determine all the necessary syntax to automatically create compilable mutants. After this MASC is designed to combine this component with the parameters to create the mutants. MASC also ensures that mutants generated for evaluation are compilable using two main steps: "(1) Using Eclipse JST's AST-based to check for identifying syntactic anomalies in the generated mutated apps, and (2) compile the mutated app automatically using build/test scripts provided with the original app." Since the portion is fully automated this is how MASC has made it possible to generate many mutants to evaluate a crypto-detector with little effort from the user.

3. Identifying Target Locations and Seeding Mutants: To locate target locations to seed mutants using the similarity scope the MDroid+ mutation analysis project was leveraged as a component of MASC. For the original work the process it used to determine mutant locations was changed to fit the scope of MASC and support was added to include dependencies that crypto mutations introduce. When I began work on MASC this component was connected to MASC where it needed to be but contained a lot of unnecessary parts that MASC was not utilizing. For the extension I identified the parts of MDroid+ that MASC was using and fully integrated them within MASC. This was done to help consolidate MASC into one project rather than a project using components from two additional projects. The components that were leveraged from MDroid+ still exist but now are fully integrated within MASC. Similarly MASC also extended μ SE to create the exhaustive scope. The extended μ SE was used to find locations where crypto-APIs can be inserted in a program and still allow the program to be compilable. Just like MDroid+ I integrated the parts of μ SE that MASC utilized into the main program. In addition the same flaw of corner cases with mutants causing compilation errors still exists within MASC. Due to how MASC is implemented this is something that is not easy to change. These cases still have a chance to appear in 0.098

3.0.1 New Features

While many of the new additions to MASC came in the form of extending work that was already there and heavily extended the evaluation. I also introduced a couple of new features into MASC to help make the framework more user friendly and allow a new perspective of evaluation to be conducted.

3.0.2 Sensitivity Evaluator

The sensitivity evaluator is a new component of the MASC framework designed as an alternative way for users evaluate crypto-decorders. This tool is designed to help users access the operators without having to understand the specifics of what each operator does. In security analysis there are sensitivities that are commonly discussed in relation to static analysis tools. These tools are built with these sensitivities in mind and claim to be able to logically handle some of them. The main sensitivities I found when look through past work were flow sensitivity, alias sensitivity, context sensitivity, path sensitivity, and object sensitivity. For this work the sensitivities are defined as the following:

Flow Sensitivity - Flow sensitivity is an incredibly precise form of sensitivity. It recognizes the order that statements are performed and can keep track of the state of the program at that point in time. A flow sensitive analysis performs its analysis based on the sequence of statements. It can tell if two variables are assigned after line 23 while a flow insensitive analysis will only know that the two variables were assigned at some point within the scope of their analysis. A flow analysis will only take into consideration portions of the program that would be run based on the previous lines. Flow sensitivity analysis is a extremely expensive computationally.

Alias Sensitivity - Alias sensitivity is typically a variation of context or flow sensitivity. In Java alias sensitivity is typically type based, this is because Java is a type safe language. Alias sensitivity is the ability to keep track of a variable that has been aliased to another

variable and still keeping track of the value. If there is a variable named `x` that equals 1 and we pass this into a method this passed value would be an alias of `x`.

Context Sensitivity - Context sensitivity takes into account the information throughout the program when method calls are made to determine if there is a vulnerability. It can differentiate between two different function calls to the same method with different variables. A context insensitive approach would flag both function calls if one of them was considered vulnerable while a context sensitive approach can differentiate between the two calls. A less sophisticated version of this is interprocedural sensitivity.

Path/Conditional Sensitivity - Path sensitivity only takes into consideration paths through the program that are feasible. It has a heavy focus on things such as conditionals. Within programs some paths or statements can not be reached by the code, a path sensitive analysis would not flag a vulnerability that is unreachable. Path sensitive analysis is only concerned with the path of the program that is possible to be executed.

Field/Index Sensitivity - Field sensitivity is the ability to differentiate different fields that are a part of the same object. If an object contains two variables one tainted and one that is not tainted and the non tainted one is called a field sensitive analysis would not flag the object as a vulnerability. This requires keeping track of all the contents of objects separately and understanding when certain aspects of an object are called by the program.

Object Sensitivity - Object sensitivity takes into account different versions of the same object. It has the ability to understand the difference between a version of an object that contains a vulnerability and one that does not contain a vulnerability. If we create two versions of object `FOO` called `f1` and `f2` and place a vulnerability in `f2` but only interact with `f1`, an object sensitive tool would be able to recognize that there is not vulnerability taking place.

With these definitions in mind for the sensitivities I designed the sensitivity evaluator. This tool allows MASC to be run and generate mutants that fit into under each specified definition. Once these definitions were clearly defined, I categorized each operator into the category or categories that it fit under. Then with this knowledge I designed the sensitivity

runner so that MASC would only produce output based on a specified sensitivity. This was done to lower the barrier of entry for MASC as well as create a new way of evaluating crypto-detectors using MASC. By defining operators in this way it is possible to put more emphasis on cases that fit the description of a crypto detector. If a crypto detector made a claim that it was flow sensitive now it is possible to easily run all the cases that are defined to be flow sensitive and see how well it performs against those mutants. It would be expected that a crypto detector that claims to handle a certain sensitivity would be better suited to handle those cases.

The sensitivity evaluator combines the input of all of the various cases and allows the user to specify how they want to run it in one file. It then handles creating all the operators that are related to the selected sensitivity. It will create the operators with the parameters provided by the user within each operator. This tool should make MASC more accessible to those with some knowledge of security without having to get a full grasp about how the operators and specifics of MASC work. I believe both the user interactivity aspect of this and the new perspective will be greatly beneficial to performing the research.

3.0.3 Automated Evaluation

In the original paper all analysis from all the crypto detectors was done by hand so it required many man hours and double checking to ensure there were no errors. To help determine results for researchers and users an additional tool was created for MASC. This tool is the automated analysis. This allows researchers to run MASC on certain crypto detectors and MASC will automatically parse the results and let the user know if the crypto detector failed and how it failed. This tool can also take output from various crypto detectors that were given MASC code and can tell where the crypto detector failed. This is done using a SARIF parser that was created. SARIF stands for Static Analysis Results Interchange Format. This file format is becoming the standard output for Static Analysis tools and is being pushed heavily by GitHub. At the time of creating this tool this file format was fairly new and did not have many tools created to parse its contents. To

build the automated analysis component it required first to build a tool that could parse the SARIF format. SARIF format shares a lot of similarities with JSON so I leveraged some Java JSON libraries such as the Google JSON simple library. I built a SARIF layer on top of this library to create a tool for parsing output. The SARIF parser tool I created takes in two SARIF files as input one of the code before mutation that was passed into the crypto detector and one Sarif file that was mutated and seeded with misuses. This is done so that if there were any misuses present in the program before MASC was run that are not taken into account and added to the total of misuses the crypto detector found because of MASC.

More specifically, I created a tool that is able to parse a SARIF output from a Crypto-API Misuse Detector to determine which seeded misuse cases were caught by the detector. Put another way, this component checks to make sure the Crypto-API Misuse Detectors were actually able to catch the misuses seeded in the program. Currently, this tool only works for the Main scope of the MASC Framework. However, it can easily be expanded to work with the Exhaustive and Selective scopes in the next extensions of the MASC Framework.

In terms of how the tool works, it takes the SARIF output files obtained from the Crypto-API Misuse Detectors as input along with the MASC properties file used for mutation. Using this information the tool parses through the MASC properties file to determine where the created mutated Java files were placed and what the Crypto-API name is being tested. Using this information, the mutated Java files that were created are scanned to find the line that contains the misuse. Once the misuses are found it is then possible to scan through the results of the SARIF file using Flow Analysis.

Flow Analysis in terms of how it is used in this project is the idea that each misuse created by an operator can be represented as a different level of complexity. MASC inherently has levels of complexity already designed into it. To put it simply, a Crypto-API Misuse Detector will have an easier time catching a misuse like “AES” than “A~ ES”.replace(“~”,”) since the former example is less syntactically complex than the latter example. By design-

ing it this way, the SARIF files can then be analyzed in order of complexity starting with the most basic misuse (or “base case” misuse) and increasing to the more complex ones until the Crypto-API Misuse Detector fails to catch one. If the crypto-detector fails then the analysis can stop because if a tool cannot find a less complex misuse it would not be expected to find a more complex or mutated misuse case. Implementing the evaluator in this way makes it a lot easier to determine where Crypto-API Misuse Detectors fail and saves a lot of manual effort since this can all be done in an automated fashion. This process also reduces the risk created by manual evaluation.

The tool was later expanded and used a part of the main MASC properties file. It was made possible that when MASC was run it could run the output it creates against a Crypto Detector check that output and let the user know which mutants were found. This additional step was built on top of the SARIF Parsing tool to create the full end to end automated analysis. This step was the integration of the SARIF Parsing analysis tool with the entire main project since the original tool looked at outputs of crypto detectors that had MASC mutants run on them. The new full automated analysis takes that work and brings together with MASC being run as well combines crypto-detectors such as CogniCrypt by running it as a part of the analysis.

Both the SARIF parsing tool and the combination with automated analysis should help users as well as researchers save a lot of time when examining crypto-detectors. This automated analysis is less error prone and can help automatically determine if a mutation was caught but also what the likely cause of the failure was. It can determine with its Flow Analysis what level of complexity caused the failure and help more easily identify the potential flaws found within detectors. Overall the tool helps MASC become a more complete project and continues to make it more accessible. This helps to further MASC’s goal of improving crypto-detectors by better identifying exactly where they fail.

3.0.4 MASC Web

In another attempt to make MASC user friendly. I designed the initial version of a website for MASC. Currently a second version of MASC Web is in active development using Django instead of Flask (which was originally used). The goal of the website was to introduce users to MASC and allow them to mutate files directly on the website. I integrated MASC into the website and made it so users could upload a file and specify the parameters they wanted mutated as well as the operator they wanted to use. The website would then provide them with both the original version and the mutated version of the file. Eventually the goal is to deploy the website to help increase visibility of MASC and allow for additional options for users to access MASC.

Chapter 4

Evaluation and Methodology

The goal of the evaluation was to find flaws in crypto-dectors using MASC once again. The main objectives were to (1) measure the effectiveness of the new operators at uncovering flaws in crypto-dectors and the old ones at uncovering flaws in new crypto-dectors (since the original operators were proven to be effective), (2) continue to learn the characteristics of the flaws found and their real world impacts, and (3) determine how likely flaws are to reappear in future versions of crypto-dectors. Based on these goals to evaluate MASC the same three original research questions were used to conduct evaluation with the addition of another: RQ1: Can MASC discover new flaws in crypto-dectors? RQ2: What are the characteristics of these flaws? RQ3: What is the impact of the flaws on the effectiveness of crypto-dectors in practice? RQ4: How likely is it for flaws to reappear in newer versions of crypto-dectors? To answer RQ1-RQ4, I used the same methodology described in the initial MASC paper. I took the original set of 9 crypto-dectors and found their most up-to-date versions if possible (Xanitizer is no longer available) and then added an additional 5 crypto-dectors to the set that were not evaluated in the original paper. The crypto-dectors that were evaluated were CogniCrypt, SpotBugs with FindSecBugs, QARK, LGTM, GitHub Code Security (two versions), ShiftLeft Scan, Amazon Code Guru, SonarQube, Codiga, Deepsource, and Snyk. All tools used are under active maintenance since the goal of MASC is to provide feedback on potential improvements to the crypto-dectors it is important that

they are still in active use. Similar to the original paper the results of the evaluation are not intended to demonstrate comparative advantages of tools and are not intended as an endorsement of any such tools, each tool is evaluated separately from the others using the same techniques. For evaluation of the crypto-detectors the steps used to conduct evaluation remained the same. I expanded on the techniques used but ensured they were still consistent with the original work since RQ1-RQ3 are the same as the original paper.

Step 1 - Selecting and mutating apps: In the original work 13 open source Android apps and four sub-systems of Apache Qpid Broker-J were found and used for mutation. For the extension these same apps with their original mutations were used as a part of the evaluation. In addition, I located 15 Android apps on Github and Tink, a Java project from Google, for mutation. To locate these additional candidates for mutation I used a similar methodology as the original work. The way the new Android apps were located was by using GitHub's advanced search to locate Android apps that had a least 200 stars, these were then sorted by how recently they were updated to help ensure they would be both compilable and up to date with dependencies. All apps that were used were tested to ensure they were compilable prior to mutation. In addition to help get a variety of different types of apps I also searched for apps with that contained tags such as: security, cryptography, secure, games, tools, calendar, exercise, and location. Out of the 15 new Android apps 5 of them were specifically found to include crypto-graphic API's (Cipher, MessageDigest, X509TrustManager) and were mutated using the Similarity scope. In addition to the 20,303 mutations found in the original MASC evaluation, I introduced 26,765 mutants within the new Android apps and Java project. This gives a grand total of 47,068 mutants generated by MASC used for evaluation, more than double the original amount of mutants. Generating all the new mutants took about 20 minutes total on MASC which addresses RQ3, and did not require any human intervention. MASC is capable of generating mutants very quickly the bottleneck of time comes from crypto-detectors evaluating the mutants.

Step 2 - Evaluating crypto-detectors and indentifying unkilld/undetected mutants: To conduct the evaluation on a crypto-detector I analyzed the mutants that were produced using the

crypto-detector. After this mutants that were not detected by the crypto-detector were indentified and flagged as unkilld. To conduct this full evaluation I kept track of logs produced by MASC for all the apps, this log contained information such as line numbers where mutations were placed and what type of mutations were placed in those locations. Then intially I ran the unmutated version of the app against the crypto-detector. Using the output produced I made a not of any crypto-graphic related misuses that were found on the unmutated version to ensure that these were not counted as a mutation produced by MASC. Then the mutated version of the app could be analyzed by the crypto-detector and results could be compared. Once the similarities of the reports were eliminated I looked at the remaining results and any remaining cryptographic misuse that was found is considered a mutant inserted by MASC that was killed by the crypto-detector. Each killed mutant is confirmed by comparing it with the log to ensure that a mutant was inserted at this location and the correct flag was produced by the crypto-detector. In addition, this analysis was also conducted automatically by the automated analysis mentioned in section 3.0.3 for any tool that could produced SARIF output such as CogniCrypt. The approach taken for this is done to ensure that misuses found by the crypto-detector that were not inserted by MASC are not considered as part of the evaluation. The main goal is to ensure all mutants that were inserted by MASC are located and marked as killed or unkilld. [Insert average number of undetected mutants across all detectors]

Step 3 - Indentifying Flaws (RQ1): For the apps using the exhaustive approach I randomly analyzed many of mutants that went undetected to locate porential flaws. I took the same approach as the original work and looked at misuses that the crypto-detector claims is can detect but in reality fails to detect it. To make sure all flaws were novel I exempted the exceptions that were stated specifically in the crypto-detector's documentation to ensure a fair evalaluation. After all the goal is to determine flaws where crypto-detectors make claims but fail to meet those claims. Since this work was done in the original MASC the claim remains consistent that "while a crypto-detector may seem flawed because it does not detect a newer, more recently identified misuse pattern, we confirm that all the

flaws we report are due to misuse cases that are older than the tools in which we find them." This is because misuses that were used to evaluate the detectors are still the most discussed misuses between 1998-2022 and since all the tools used either have new versions or were recently deployed (Amazon Code Guru). In addition to confirm flaws I used the minimal app created for the original work that contained only the undetected misuses to re-analyze the all the tools. This app was also altered to include mutations created by the new operators to verify these as well. Step 4 - Characterizing flaws (RQ2): Flaws were grouped into the same flaw classes, based on most likely cause, used for the original evaluation. The only difference was the flaw classes were updated to include the new mutants but the original flaw classes still remain intact. This was done to help further grasp why are the tools failing. In addition all 13 crypto-detectors were evaluated using the minimal examples represented by each flaw to gain a further understanding as to why a flaw might be missed despite the documentation make claims as such. Flaws were reported to their respective tool maintainers, and so far SonarQube has created 5 high priority issues in response to the reported flaws. Step 5 - Understanding the practical impact of flaws (RQ3): To answer this research question for the extension I analyzed repositories to ensure that the new misuses that were introduced could also be found in real world applications. This was done using GitHub Code Search and was additionally confirmed manually. Step 6 - Attributing flaws to mutation vs base instantiation: To ensure that a flaw that MASC found was due to a mutation and not simply a lack of coverage of the base case, each crypto-detector was evaluated with the most basic version of the misuse to determine if it was able to catch that. An example of this would be directly testing something such as `Cipher.getInstance("DES")`, if this use case is caught it is then possible for it to detect mutations of this statement. Step 7 Reevaluation of the original 8 crypto-detectors: To reevaluate the original crypto-detectors I used the same approach that was conducted on them in the original work. I ensured that they had new versions since the original work was published. To ensure that what I found was a flaw that reappeared rather than a flaw that had not been fixed from the original work, I made sure that the issue that was reported

by the original authors was closed. This was to ensure that the bug was addressed. I also report if any of the flaws present in the original work that were never fixed are still present in the new version.s

Chapter 5

Results and Findings

5.0.1 Analyzing Flaws

The new analysis of the 13 crypto detectors results in [work is currently ongoing] new flaws being found. Additionally the analysis found [insert number of flaws] in the 5 new crypto-detectors that were located in the original flaws found by MASC. For this work I included an updated version of the flaw classes with the newly located flaws and present an updated table that maps the flaws to the new versions of the crypto-detectors and new crypto-detectors. Similarly, to the original work I found that the majority of total flaws found could be attributed to mutations [insert values] while [insert value] could be found by just using the base case versions of the misuse. This continues to prove that mutation testing provides better results for finding flaws than just using the base cases alone to determine flaws. I also found that [insert number] of flaws found in the original paper are still present in the reevaluated new versions of the crypto-detectors. This shows that flaws do have a tendency to reappear and demonstrates that evaluation of crypto-detectors is not a one time thing, it needs to be ongoing. Each release has the potential to reveal more bugs and possibly bring back old misuses. In addition across the reported flaws a total of [insert number] are still present in the new versions of the crypto-detectors even after these issues were reported. Since I reevaluated crypto-detectors found in the original paper I used the patched versions that included a fix for the multidex issue presented in the initial paper.

However I noticed, notably with Amazon Code Guru, there does appear to be a limit to the number of recommendations it creates for a specific misuse. For example, when used the exhaustive scope the misuse produced a "weak cipher algorithm" warning but it only reported 100 of these despite many more being seeded within each program. This would then mean after the user fixes the flaws they would have to know to run another analysis on their code. Similar to the first group this does affect the reliability of the results and makes it challenging to determine how well it evaluates many misuses. I will further discuss this fault with Amazon Code Guru and the crypto-dectors that did not fully analyze the code in the Discussion section [add section id]. These types of flaws are included in Flaw Class Zero since they cannot directly be attributed to being found by MASC.

NOTE: Work is still ongoing for this portion. I am still waiting on a response from some of the tool maintainers. In addition some evaluation is still be conducted and results are still being processed. This will be updated prior to the final submission. Since work is still ongoing the new flaws are still being identified and are not yet included in the flaw class table.

FC1: String Case Mishandling (F1): This class was the motivating example found for the original work and for consistency remains in its own class for this work. This example can be found in Table 1 as F1. For this case a developer may use des or dEs (instead of the expected DES) in `Cipher.getInstance(<parameter>)` without any errors being raised. This was not found by Snyk despite being able to detect the base case version of this. Condiga and QARK were also unable to detect this flaw but were unable to detect this base case as well. FC2: Incorrect Value Resolution (F2 – F9): For these flaws 11/13 crypto-dectors failed to successfully detect all 9 flaws. Notably though the crypto-dectors that were evaluated in the initial paper did perform better on average than the 5 new tools that were evaluated. Snyk and SonarQube performed on par with some of the tools that were initially evaluated. The crypto-dectors that had flaws reappear in this class are ... SonarQube was able to detect some of the more simpler cases but failed due to the tool not evaluating method invocations. SonarQube is only capable of detecting String literals and

can detect it if it is contained in a variable. For the rest of this section we are still waiting to hear back from tool maintainers to add the reasoning as to why their flaws were not evaluated by the tool. FC3: Incorrect Resolution of Complex Inheritance and Anonymous Objects (F10 – F13): Flaws in this class occur because [insert value of crypto-dectors] were unable to resolve the complex inheritance relationships among classes. These flaws were found specifically by using the flexible mutation operators. As found in the original paper this is clearly a consideration for some of the crypto-dectors while others did not consider this in design.

FC4: Insufficient Analysis of Generic Conditions in Extensible Crypto-APIs (F14 – F16): The flaws in this section represent the inability of crypto-detectors to identify fake conditions and also true conditionals. This determines if a crypto-detector is capable of following path sensitivity.

FC5: Insufficient Analysis of Context-specific Conditions in Extensible Crypto-APIs (F17 – F19): The flaws found within this class are similar to those found in FC4, however, the fake conditionals are contextualized to the overridden function. This would simulate the behavior of an evasive developer because one may try to add further realism to fake conditions to avoid tools that are capable of detecting simple generic conditions.

5.0.2 Exhaustive Results

So far results have been found for the the 5 new operators. However, results for this section are still ongoing. The apps are currently being run on the original crypto-dectors. Results will be reported by the final version.

Chapter 6

Conclusion

MASC has been demonstrated once again as being effective at finding flaws in crypto-detectors. The work has been greatly expanded but there are still more areas for further improvement. This new extension of MASC has also reintroduced some areas of discussion initially brought to light by the original work. Since some time has passed some of the perspectives of these discussions have shifted.

6.0.1 Limitations

MASC is still designed to help find flaws in crypto-detectors and still cannot guarantee all flaws in a crypto-detector will be found. In fact formal verification should still be done in conjunction with MASC. MASC's design from the very beginning was to allow for "systematic evaluation of crypto-detectors, which is an advancement over manually curated benchmarks." This still remains true while the extension has expanded MASC's functionality and coverage it is still not representative of all misuses. MASC is still held back by the following limitations: 1. Completeness of the Taxonomy: The same approach was used to ensure the taxonomy was comprehensive as the original work. All steps were performed carefully and utilized the same best practices found in other work. However, it is still possible that some cases or sources could be missed during extraction. The extra evaluation step was added to help ensure accuracy but it is still possible that some subtle

contexts were once again missed. Up to current date though I believe that this is still, to the best of my knowledge, the most comprehensive found in recent works taxonomy in this space.

2. Focus on Generic Mutation Operators: This was a concern of the original paper since the goal was to apply as many misuses from the taxonomy as possible. This was a main area of expansion for this extension. MASC now contains many more new operators that do represent some of the more specific cases found in the taxonomy. However, the taxonomy is still not fully represented and many of the new operators were still designed to have multiple use cases. Priorities for new operators were still focused on covering more potential misuse even though the focus was exclusively on expanding the restrictive operators.
3. Focus on Java and JCA: MASC's approach is still informed by JCA and Java. No work has been done since the original paper to adapt MASC to other languages at this time.
4. Evolution of APIs: Since changes are still made to how JCA operates this may eventual lead to changes being necessary in MASC. Up to now MASC is current and the operators from the original paper do still function correctly. However, as time passes some changes may become necessary despite MASC using reflection and automated code generation to ensure flexibility. In addition it is still an ongoing project to include more misuse cases within MASC to get closer to fully representing the current taxonomy and beyond.
5. Relative Effectiveness of Individual Operators: My research did not look into this limitation that was present in the original paper. This paper looks further into what MASC is capable of doing as a whole but does not evaluate how effective each operator is at finding flaws. This still requires its own study.
6. Consistency with the Original Work: I was in constant communication with the original authors and confirming with them how the original work was created. I incorporated feedback from them and had them review changes to the tool to help ensure as much consistency as possible. In addition for evaluation and expansion I followed all the steps laid out by the original work to ensure that the work could stand as one whole project rather than an add on. I believe that this was done to the best of my ability and MASC now stands as an expanded and more

comprehensive tool. Since I was not directly involved in the original work there may be some small details that were not done exactly the same as the original work.

6.0.2 Discussion

1. Security-centric Evaluation Design MASC still places a heavy focus on security centric design. From the perspective of this framework it is believed that security should come first no matter how unlikely or evasive a case may be. This is part of why some of the new operators were designed to create more evasive cases. This idea still clashes with the idea that designers of tools look more into a technique-centric perspective. Many of these tools are not designed with a threat model in mind or directly from a security prespective. In fact many of the tools analyze best coding practices in addition to looking for security misuses. This leads to a gap between what is expected of a tool that claims it can perform security checks. 2. What is scope for Technique-centric Design There is still an ongoing discussion on the space on what the scope should be. Should there be more of a focus placed on what is most common and oncommon or should it be on what can easily be computed statically vs cannot easily be computed. Even from a security prespective this is challenge since common misuses would be exptected to be caught, however, new misuses pose a bigger threat due to the lack of awareness of the looming threat. Optimially most tools would like to cover all possible cases but since this is not feasible the debate is what is most important and what should be expected on crypto-detectors. 3. The Need to Strengthen Crypto-Detectors Many detectors make claims with the assurance that they can detect certain misuses. When tested many of these claims are proven to still fall short of expectations. If crypto-dectors claim to be able to secure your code it is important that they actually can. This means detecting uncommon cases and being held accountable for falling short. MASC has shown that it is possible to find gaps in crypto-detectors and help them improve but they are still a long way off from being able to detect hard to statically compute misuses cases. This leads to the question of should crypto-detectors be able to make the claims they do to secure your code? 4. The original work showed that there is

interest in developers to make tools more secure. This was something they strived for and it was proven that some of the tools that were evaluated did improve since the first paper. This demonstrates that not only did the express interest they demonstrated interest. Even in the new tools that were evaluated when issues were reported to them they were eager to make a fix to improve their tool. Many developers have expectations that these tools help make their code secure and a lot of times are not aware of common misuses. The further these tools can expand with security centric design in mind the more developers can expect to rely on them.

5. Pushing toward SARIF As mentioned SARIF is a relatively new SAST output format. Some crypto-detector designers have been reached out to and expressed interest in being able to produce this format for their tool. MASC integrated this in hopes of helping to encourage more crypto-detector designers to look into this format. If there is a standard output not only does it make it easier to evaluate how well tool perform it also allows for new tools to be produced to help users evaluate their applications. Having a standard output as an expectation can lead to easier evaluation and can help create a way to fully automatize the evaluation of crypto-detector. This would help users be able to easily know how reliable any tool is before they use it. In addition if this was a standard output this would also allow tools to more easily report all misuses found. As mentioned in a prior section some crypto-detectors such as Amazon Code Guru do not provide a full report for their analysis. This is likely due to UI limitations since the whole codebase is still scanned. Having an output in the style of SARIF can help provide a full report to the user since it likely would not have to be directly consumed by the user, it would be parsed and display the output that way.

6. The Need for constant evaluation As shown in the reevaluation of the original crypto-detectors some bugs or flaws have the potential to reappear in future versions of a tool. Due to the ongoing iteration of crypto-detectors it is important for them to be reevaluated. While obviously an evaluation would be necessary when a new misuse appears reevaluating the old misuses is equally as important. Bugs have a tendency to reappear while refactoring and making changes and as found the same

thing goes for flaws. Having a constant evaluation will help ensure that once a flaw is found and eliminated that it stays that way.

6.0.3 Lessons Learned

6.0.4 Conclusion

The creation of a tool like MASC has helped shift the mindset of the crypto-detectors designers into looking further into a more security-focused design. They have made claims and as found in this extension they have improved in some areas. There is still a long way to go and by reaching out to the designers and reporting flaws I have seen that there is interest in ensuring security. With the MASC framework existing and being known to the crypto-detector makers it is possible for them to easily perform self evaluation and continue to improve. As MASC continues to expand in the future it is possible to help push crypto-detectors into become more secure as well. With MASC existing in the same ecosystem as crypto-detectors it is possible to help push towards a more security centric design for tools and one day users can truly expect when their code is analyzed that it is secure.

Appendix A

Appendix A

A.1 Code Snippets

Listing A.1: Method Chaining (**OP₅**).

```
1    Class T { String algo="AES/CBC/PKCS5Padding";
2    T mthd1(){ algo = "AES"; return this;} T mthd2(){ algo="DES";
        return this;} }
3    Cipher.getInstance(new T().mthd1().mthd2());
```

Listing A.2: Predictable/Non-Random Derivation of Value (**OP₆**)

```
1    val = new Date(System.currentTimeMillis()).toString();
2    new IvParameterSpec(val.getBytes(),0,8);}
```

Listing A.3: Exception in an *always-false* condition block (**OP₇**).

```
1    void checkServerTrusted(X509Certificate[] x, String s)
2    throws CertificateException {
3    if (!(null != s && s.equalsIgnoreCase("RSA"))) {
4        throw new CertificateException("not RSA");}
```

Listing A.4: False return within an *always true* condition block (**OP₈**).

```
1    public boolean verify(String host, SSLSession s) {
2    if(true || s.getCipherSuite().length()>=0)}
```

```
3         return true;} return false;}
```

Listing A.5: Implementing an Interface with no overridden methods.

```
1 interface ITM extends X509TrustManager { }
2 abstract class ATM implements X509TrustManager { }
```

Listing A.6: Inner class object from Abstract type (OP_{12})

```
1 new HostnameVerifier(){
2     public boolean verify(String h, SSLSession s) {
3         return true; } };
```

Listing A.7: Anonymous Inner Class Object of X509ExtendedTrustManager (??)

```
1 new X509ExtendedTrustManager(){
2     public void checkClientTrusted(X509Certificate[] chain, String a)
3         throws CertificateException {}
4     public void checkServerTrusted(X509Certificate[] chain, String
5         authType)throws CertificateException {}
6     public X509Certificate[] getAcceptedIssuers() {return null;} ...};
```

Listing A.8: Specific Condition in checkServerTrusted method (??)

```
1 void checkServerTrusted(X509Certificate[] certs, String s)
2     throws CertificateException {
3     if (!(null != s || s.equalsIgnoreCase("RSA") || certs.length >=
4         314)) {
5         throw new CertificateException("Error");}}
```

Listing A.9: Anonymous Inner Class Object of An Empty Abstract Class that implements HostnameVerifier

```
1 abstract class AHV implements HostnameVerifier{} new AHV(){
2     public boolean verify(String h, SSLSession s)
3         return true;}};
```

Listing A.10: Anonymous inner class object with a vulnerable checkServerTrusted method (F13)

```

1      abstract class AbstractTM implements X509TrustManager{} new
      AbstractTM(){
2      public void checkServerTrusted(X509Certificate[] chain, String
      authType) throws CertificateException {}
3      public X509Certificate[] getAcceptedIssuers() {return null;}}};

```

Listing A.11: Anonymous Inner Class Object of an Interface that extends HostnameVerifier

```

1      interface IHV extends HostnameVerifier{} new IHV(){
2      public boolean verify(String h, SSLSession s) return true;}};

```

Listing A.12: Misuse case requiring a trivial new operator

```

1      KeyGenerator keyGen = KeyGenerator.getInstance("AES");
2      keyGen.init(128); SecretKey secretKey=keyGen.generateKey();

```

Listing A.13: CryptoGuard's code ignoring names with "android"

```

1      if (!className.contains("android."))
2      classNames.add(className.substring(1, className.length() - 1));
      return classNames;

```

Listing A.14: Generic Conditions in checkServerTrusted

```

1      if(!(true || arg0==null || arg1==null)) {
2      throw new CertificateException();}

```

Listing A.15: Transformation String formation in Apache Druid similar to **F2** which uses AES in CBC mode with PKCS5Padding, a configuration that is known to be a misuse [? ?].

```

1      this.name = name == null ? "AES" : name;
2      this.mode = mode == null ? "CBC" : mode;
3      this.pad = pad == null ? "PKCS5Padding" : pad;
4      this.string = StringUtils.format(
5      "%

```

Listing A.16: Iterative Method Chaining

```
1      Class T {
2          int i = 0;
3          cipher = "AES/GCM/NoPadding";
4          public void A(){
5              cipher = "AES/GCM/NoPadding";
6          }
7          public void B(){
8              cipher = "AES/GCM/NoPadding";
9          }
10         public void C(){
11             cipher = "AES/GCM/NoPadding";
12         }
13         public void D(){
14             cipher = "AES";
15         }
16         public String getVal(){
17             return cipher
18         } }
19
20
21 Cipher.getInstance(new T().A().B().C().D().getVal() ) ;
```

Listing A.17: Iterative Conditionals

```
1      Class T {
2          int i = 0;
3          cipher = "AES/GCM/NoPadding";
4          public void A(){
5              if ( i == 0){
6                  if (i == 0){
7                      if(i == 0){
8                          cipher = "AES";
9                      }
10                     else{
```

```

11             cipher = "AES/GCM/NoPadding";
12         }
13     }
14     else{
15         cipher = "AES/GCM/NoPadding";
16     }
17 } else{
18     cipher = "AES/GCM/NoPadding";
19 }
20 }}
21
22 public String getVal(){
23     return cipher
24 } }
25
26
27 Cipher.getInstance(new T().A().getVal() ) ;

```

Listing A.18: Iterative Conditionals

```

1     Class T {
2         int i = 0;
3         cipher = "AES/GCM/NoPadding";
4         public void A(){
5             if ( i == 0){
6                 if (i == 0){
7                     if(i == 0){
8                         cipher = "AES";
9                     }
10                    else{
11                        cipher = "AES/GCM/NoPadding";
12                    }
13                }
14            } else{
15                cipher = "AES/GCM/NoPadding";

```



```

16         }
17     } else{
18         cipher = "AES/GCM/NoPadding";
19
20     }}
21
22     public String getVal(){
23         return cipher
24     } }
25
26
27     Cipher.getInstance(new T().A().getVal() ) ;

```

Listing A.19: Method Builder

```

1     Class T {
2         int i = 0;
3         cipher = "AES/GCM/NoPadding";
4         public String A(){
5             return "D";
6         }
7         public String B(){
8             return "E";
9         }
10        public String C(){
11            return "S";
12        }
13        public void add(){
14            cipher = A() + B() + C();
15        }
16        public String getVal(){
17            return cipher
18        } }
19
20

```

```
21 Cipher.getInstance(new T().add().getVal() );
```

Listing A.20: Iterative Conditionals, using the object created in Listing A.1

```
1 T secure = new T();
2 T insecure = new T().mthd2();
3 secure = insecure;
4 Cipher.getInstance(secure.getVal());
```

Listing A.21: Build Variable

```
1 String cryptoVariable = "AES";
2 char[] cryptoVariable1 = cryptoVariable.toCharArray();
3 javax.crypto.Cipher.getInstance(String.valueOf(cryptoVariable1));
```

Listing A.22: Substring

```
1 javax.crypto.Cipher.getInstance("secureParamAES".substring(11));
```

Listing A.23: Static Keystore

```
1 byte[] cryptoTemp = "12345678".getBytes();
2 javax.crypto.spec.IvParameterSpec ivSpec = new javax.crypto.spec.
  IvParameterSpec.getInstance(cryptoTemp, "AES");
```
